

## Highlights

### **Log-Weighted Cross-Entropy Framework for Imbalanced Classification**

Seung Ho Go, Hyunseok Lee, Jin Hee Yoon

- A logarithmic re-weighting loss that bounds the weight explosion of WCE.
- A single-parameter generalization tuned automatically on a small proxy subset.
- An  $\epsilon$ -smoothed variant keeps per-class weights bounded even for the rarest classes.
- Evaluation spans image, text, tabular, and network-intrusion domains.
- The proposed losses improve as imbalance grows and avoid collapse of rival losses.

# Log-Weighted Cross-Entropy Framework for Imbalanced Classification

Seung Ho Go<sup>a</sup>, Hyunseok Lee<sup>b</sup> and Jin Hee Yoon<sup>c,\*</sup>

<sup>a</sup>Sejong University, Department of Software 209 Neungdong-ro, Gwangjin-gu, Seoul, Republic of Korea

<sup>b</sup>Sejong University, Department of Data Science 209 Neungdong-ro, Gwangjin-gu, Seoul, Republic of Korea

<sup>c</sup>Sejong University, Department of Artificial Intelligence and Information Technology 209 Neungdong-ro, Gwangjin-gu, Seoul, Republic of Korea

## ARTICLE INFO

### Keywords:

Class imbalance  
Long-tailed learning  
Cross-entropy loss  
Cost-sensitive learning  
Weight smoothing  
Network intrusion detection

## ABSTRACT

Class imbalance is a pervasive obstacle in machine learning, arising whenever a few majority categories dominate the training data while critically important minority categories remain severely underrepresented. Among algorithm-level remedies, Weighted Cross-Entropy (WCE) is the most direct: it scales the loss of each class by the inverse of its frequency. However, these inverse-frequency weights grow without bound as the imbalance becomes more severe, producing unstable optimization. More elaborate alternatives such as Class-Balanced (CB) Loss and Focal Loss introduce additional hyperparameters and, in our experiments, behave inconsistently across datasets—occasionally falling below ordinary Cross-Entropy (CE). We propose Log-Weighted Cross-Entropy (LWCE), which replaces the inverse-frequency weight with a logarithmically compressed counterpart that keeps a meaningful boost for rare classes while bounding weight growth, and which introduces no additional hyperparameters. We further generalize it to Power-Log Weighted Cross-Entropy (PLWCE) through a single exponent that adjusts the focusing strength and is selected automatically on a small proxy subset, and we introduce a complementary  $\epsilon$ -smoothed variant (ES-LWCE) that adds a small constant to the logarithmic denominator, providing an explicit, user-controlled cap on the per-class weight and keeping it well defined even for vanishing class counts, such as classes absent from a mini-batch. To assess generality, we evaluate the proposed losses across four distinct application domains that together span a wide range of imbalance severity and model architectures: long-tailed image classification, tabular classification benchmarks, network intrusion detection, and text classification with a Transformer encoder. Across these domains, LWCE and PLWCE improve over Cross-Entropy as the imbalance grows and, unlike the competing re-weighting losses, avoid the severe degradation below plain Cross-Entropy that these baselines occasionally exhibit. These results indicate that logarithmic weight compression offers a simple, stable, and parameter-efficient alternative to existing re-weighting strategies for class-imbalanced classification.

## 1. Introduction

Across a wide range of real-world classification tasks—including medical diagnosis, financial fraud detection, network intrusion detection, and visual recognition—the categories of interest are rarely observed in equal proportions. Instead, a small number of majority categories concentrate the overwhelming majority of the available samples, while many minority categories are represented by only a handful of examples (He and Garcia, 2009; Johnson and Khoshgoftaar, 2019). This long-tailed structure is not an artifact of poor data collection but an intrinsic property of the underlying phenomena: healthy patients vastly outnumber those with a rare disease, legitimate transactions dwarf fraudulent ones, and benign network traffic far exceeds malicious traffic. Because the minority categories are often precisely the ones that carry the greatest practical cost when missed, building classifiers that remain reliable under imbalance is a problem of broad and lasting importance.

The central difficulty is that the standard training objective is itself biased toward the majority. When a model is trained with ordinary Cross-Entropy (CE) loss, the aggregate gradient signal is dominated by the numerically abundant classes, so the optimizer can drive the overall loss down while largely ignoring rare categories. The resulting classifier attains a high overall accuracy yet performs poorly on exactly the minority classes that matter most—a failure mode commonly described as majority bias. Standard accuracy therefore masks the problem rather than revealing it, and a method that addresses imbalance must improve minority-class performance without sacrificing the majority.

\*Corresponding author

 gseungho999@gmail.com (S.H. Go); hslee901@gmail.com (H. Lee); jin9135@sejong.ac.kr (J.H. Yoon)  
ORCID(s): 0009-0006-8715-9592 (S.H. Go); 0009-0002-0895-8785 (H. Lee); 0000-0002-1437-1350 (J.H. Yoon)

Existing remedies fall broadly into two families. Resampling methods reshape the training distribution directly: oversampling techniques such as the Synthetic Minority Over-sampling Technique (SMOTE) synthesize new minority instances (Chawla, Bowyer, Hall and Kegelmeyer, 2002), while undersampling discards majority examples to restore balance. Both directions carry well-known risks—oversampling can introduce synthetic artifacts and encourage overfitting, whereas undersampling discards potentially informative majority data (Drummond and Holte, 2003). Algorithm-level methods instead modify the training objective. These include margin-based losses that enlarge the decision boundary around minority classes (Cao, Wei, Gaidon, Aréchiga and Ma, 2019), logit-adjustment techniques that calibrate class priors, and loss re-weighting methods that scale the per-class loss by class-dependent weights. Among these, re-weighting is especially attractive for three practical reasons. Unlike resampling, it leaves the training data untouched and therefore introduces neither synthetic artifacts nor information loss; unlike architecture-level remedies, it adds no parameters and can be dropped into virtually any network without altering the model; and it requires only a minimal, modality-independent change to the loss function. Indeed, although many sophisticated long-tailed learning methods have since been proposed, re-weighting remains one of the few approaches that can be applied to almost any architecture without modifying the model itself. For these reasons this paper focuses on the re-weighting family.

Despite its appeal, re-weighting remains unsatisfying in practice. Weighted Cross-Entropy (WCE), the most direct re-weighting scheme, assigns each class a weight inversely proportional to its frequency,

$$w_c^{\text{WCE}} \propto \frac{1}{n_c}. \quad (1)$$

As a minority class becomes rarer and  $n_c$  shrinks, this inverse-frequency weight grows without bound relative to the majority, so the corresponding gradients become disproportionately large and destabilize optimization under extreme imbalance. More sophisticated alternatives were designed to soften this behavior but bring their own drawbacks. Focal Loss reduces the influence of easy, typically majority, examples through a tunable focusing term, but it operates on per-sample confidence rather than class frequency and therefore behaves inconsistently across imbalance levels. Class-Balanced (CB) Loss re-weights by an effective number of samples governed by an additional hyperparameter, yet that hyperparameter is sensitive and, in our experiments, CB Loss degrades below plain Cross-Entropy under severe imbalance. Existing methods thus attempt to alleviate this instability through additional mechanisms, yet none directly addresses the uncontrolled growth of the inverse-frequency weight while preserving the simplicity of WCE. In short, the field still lacks a re-weighting loss that is simultaneously simple, hyperparameter-light, numerically stable, and dependable across both the level of imbalance and the application domain.

To address this gap, we propose Log-Weighted Cross-Entropy (LWCE). The key idea is to replace the inverse-frequency weight of WCE with a logarithmically compressed counterpart, applying a logarithm to the class count before inverting it. This preserves a meaningful emphasis on rare classes—the weight still decreases with class size—but fundamentally bounds how fast that emphasis can grow, so the explosive behavior of WCE is removed. Crucially, the resulting loss introduces no additional hyperparameters and can be used as a drop-in replacement for Cross-Entropy. We further develop LWCE in two complementary directions. The first, Power-Log Weighted Cross-Entropy (PLWCE), controls *how strongly* rare classes are emphasized: it introduces a single exponent that continuously adjusts the focusing strength, ranging from near-uniform weighting to an aggressive emphasis on the tail. This exponent is the only free parameter, and it is selected automatically through a lightweight grid search on a small proxy subset, adding negligible computational overhead. The second,  $\epsilon$ -Smoothed Log-Weighted Cross-Entropy (ES-LWCE), instead adds a small constant  $\epsilon$  to the logarithmic denominator and reduces to LWCE as  $\epsilon \rightarrow 0$ . The base LWCE weight is already bounded for every class that appears at least once, so the role of ES-LWCE is not to make an otherwise unbounded weight finite, but to provide an *explicit, user-controlled* cap of  $1/\epsilon$  on the per-class weight and to keep it well defined in the degenerate case where a class is temporarily absent—for example, when class frequencies are estimated per mini-batch or maintained online for streaming or growing label sets—without changing the relative ordering of class weights. It is therefore best viewed as a robustness extension for such vanishing- or drifting-count settings rather than as a separate competitor on datasets with fixed, strictly positive counts. The two extensions are otherwise orthogonal: PLWCE adjusts the *strength* of the re-weighting, whereas ES-LWCE controls its *maximum weight* and smooths the extreme tail. Although the two can in principle be combined into a single power-log-smoothed weight, we study them separately to isolate the effect of focusing strength from that of smoothing.

We assess the proposed losses through a deliberately broad empirical study rather than a single benchmark. To test whether the benefit of logarithmic compression generalizes, we evaluate across three application domains that

together cover a wide spectrum of imbalance severity: long-tailed image classification on standard vision benchmarks, a collection of tabular classification datasets drawn from public repositories, and several network intrusion detection datasets in which rare attack types are extremely scarce. In every domain we compare the same set of re-weighting losses under a controlled protocol with repeated runs, and we report metrics that explicitly reflect minority-class performance rather than overall accuracy alone. This cross-domain design lets us examine not only whether the proposed losses help, but how their advantage changes as imbalance becomes more extreme.

The main contributions of this paper are as follows.

1. We propose Log-Weighted Cross-Entropy (LWCE), a parameter-free drop-in replacement for Weighted Cross-Entropy that substitutes the inverse-frequency weight with a logarithmically compressed one, removing the weight-explosion failure mode while retaining a meaningful emphasis on rare classes.
2. We develop two complementary extensions of LWCE: Power-Log Weighted Cross-Entropy (PLWCE), whose single exponent—tuned automatically on a small proxy subset—controls the focusing strength, and  $\epsilon$ -Smoothed Log-Weighted Cross-Entropy (ES-LWCE), whose smoothing constant imposes an explicit, user-controlled cap on the per-class weight and keeps it well defined even for vanishing or mini-batch-level class counts.
3. We analyze the theoretical properties of the proposed weighting functions, showing that their rare-to-frequent weight ratio grows only logarithmically with the imbalance ratio—in contrast to the linear growth of WCE—and that the resulting weights remain bounded and numerically stable even for the rarest classes.
4. We provide a controlled cross-domain evaluation spanning image, tabular, and network intrusion detection data, showing that the proposed losses improve over Cross-Entropy as imbalance grows and, unlike Class-Balanced Loss and Focal Loss, avoid the severe degradation below plain Cross-Entropy that these baselines occasionally exhibit.

The remainder of this paper is organized as follows. Section 2 reviews related work on learning from imbalanced and long-tailed data. Section 3 defines the proposed losses and analyzes their properties. Section 4 describes the experimental setup across the three domains and presents the results. Section 5 concludes and discusses limitations and future work.

## 2. Related Work

Methods for learning under class imbalance can be organized by the level at which the imbalance is addressed: the *data*, by reshaping the training distribution; the *loss*, by making the training objective cost-sensitive; the *representation*, by rebalancing the learned feature space or classifier; and the *architecture*, by restructuring the model itself (He and Garcia, 2009; Johnson and Khoshgoftaar, 2019; Krawczyk, 2016; Branco, Torgo and Ribeiro, 2016; Zhang, Kang, Hooi, Yan and Feng, 2023; Chen, Yang, Yu, Shi and Chen, 2024). Table 1 summarizes this taxonomy with representative methods. Our proposal is a loss-level method and, more specifically, a frequency-based re-weighting loss. We review all four levels for context and then position logarithmic re-weighting among existing cost-sensitive losses.

**Table 1**

A taxonomy of approaches to class-imbalanced learning, organized by the level at which the imbalance is addressed. The proposed losses are loss-level, frequency-based re-weighting methods.

| Level                | Representative methods                       |
|----------------------|----------------------------------------------|
| Data-level           | SMOTE, ADASYN, under-/over-sampling          |
| Loss-level           | WCE, CB Loss, Focal Loss, <b>LWCE (ours)</b> |
| Representation-level | Decoupled training, distribution alignment   |
| Architecture-level   | Multi-branch, multi-expert models            |

### 2.1. Data-level resampling

Resampling rebalances the class distribution before or during training. Oversampling synthesizes additional minority instances—the Synthetic Minority Over-sampling Technique (SMOTE) interpolates between neighbouring minority samples (Chawla et al., 2002), and ADASYN adapts the amount of synthesis to local difficulty (He, Bai,

Garcia and Li, 2008)—whereas undersampling discards majority examples and can outperform oversampling on classical learners (Drummond and Holte, 2003). Although effective and classifier-agnostic, synthetic oversampling may introduce noisy or unrealistic samples and encourage overfitting, while undersampling discards potentially informative data (Buda, Maki and Mazurowski, 2018; Branco et al., 2016). These trade-offs motivate objective-level alternatives that leave the data untouched.

## 2.2. Cost-sensitive and re-weighting losses

Algorithm-level methods instead make the loss itself imbalance-aware. We group them by the signal that drives the re-weighting into frequency-based, difficulty-based, and margin- or logit-based methods. *Frequency-based* methods scale each class loss by a function of its count. The most direct instance is Weighted Cross-Entropy, which sets  $w_c \propto 1/n_c$  so that the weight is inversely proportional to class count (Wang, Liu, Wu, Cao, Meng and Chen, 2019); these weights, however, grow without bound under extreme imbalance and destabilize optimization, and the benefit of such importance weighting is known to diminish for over-parameterized networks driven to low training error (Byrd and Lipton, 2019). Class-Balanced Loss softens this by weighting with an effective number of samples,  $w_c \propto (1 - \beta)/(1 - \beta^{n_c})$ , but this introduces an extra hyperparameter  $\beta \in [0, 1)$  (Cui, Jia, Lin, Song and Belongie, 2019). *Difficulty-based* methods modulate each example by its own hardness rather than by its class: Focal Loss down-weights easy, mostly majority, samples by a factor  $(1 - p_i)^\gamma$ , where  $p_i$  is the predicted probability of the true class and  $\gamma$  a tunable focusing exponent (Lin, Goyal, Girshick, He and Dollár, 2017), and the Gradient Harmonizing Mechanism flattens the distribution of gradient norms so that both easy samples and hard outliers are suppressed (Li, Liu and Wang, 2019). *Margin- and logit-based* methods reshape the decision geometry rather than the per-class loss magnitude: they enlarge minority margins (Cao et al., 2019), equalize gradients across frequencies (Tan, Wang, Li, Li, Ouyang, Yan and Yan, 2020), adjust logits by class priors (Menon, Jayasumana, Rawat, Jain, Veit and Kumar, 2021), or learn the weights themselves through meta-learning and cost-sensitive representation learning (Ren, Zeng, Yang and Urtasun, 2018; Ren, Yu, Sheng, Ma, Zhao, Yi and Li, 2020; Khan, Hayat, Bennamoun, Sohel and Togneri, 2018), with recent work casting weight learning as an optimal-transport problem between the imbalanced and a balanced distribution (Guo, Li, Zheng, Zhao, Zhou and Zha, 2022). These losses improve over plain Cross-Entropy in many settings, but they introduce additional, often sensitive, hyperparameters and can behave inconsistently across datasets. Our logarithmic family belongs to the frequency-based group and is therefore complementary to the difficulty- and margin-based axes; it targets the central weakness of frequency-based re-weighting by bounding the frequency weight without introducing sensitive hyperparameters.

## 2.3. Representation learning and decoupling

A complementary line of work argues that imbalance harms the classifier more than the feature extractor. Decoupling representation and classifier learning (Kang, Xie, Rohrbach, Yan, Gordo, Feng and Kalantidis, 2020) and distribution alignment (Zhang, Li, Yan, He and Sun, 2021) rebalance only at the classification stage, while open long-tailed recognition addresses imbalance jointly with open-set and few-shot effects (Liu, Miao, Zhan, Wang, Gong and Yu, 2019); recent surveys provide a broad taxonomy of these directions (Zhang et al., 2023). This concern extends to the era of large pretrained and foundation models, where the fine-tuning strategy itself interacts with the long tail and must be adapted to avoid harming the rarest classes (Shi, Wei, Zhou, Shao, Han and Li, 2024). Such methods are largely orthogonal to the choice of loss and can be combined with re-weighting, so our contribution—confined to the loss term—composes with them rather than competing.

## 2.4. Architecture-level and multi-expert methods

A third family addresses imbalance by restructuring the network rather than the data or the loss. Multi-branch designs train separate pathways for the head and tail of the distribution and fuse them through cumulative, bilateral learning (Zhou, Cui, Wei and Chen, 2020). Mixture-of-experts and multi-expert approaches extend this idea by training several diverse, distribution-aware experts and then routing or aggregating their predictions: RIDE routes inputs to complementary experts (Wang, Lian, Miao, Liu and Yu, 2021), ACE assembles ally-complementary experts in a single stage (Cai, Wang and Hwang, 2021), LFME distills knowledge from multiple experts trained on balanced subsets (Xiang, Ding and Han, 2020), and test-agnostic aggregation adapts the expert ensemble to an unknown test distribution (Zhang, Hooi, Hong and Feng, 2022). Such designs can deliver strong tail accuracy, but they increase parameter count and inference cost, and—like the resampling and representation methods above—they are orthogonal to the training objective: a re-weighting loss such as ours can be applied within each branch or expert rather than replacing the architecture.

## 2.5. Positioning of the proposed losses

Within the re-weighting family, existing schemes face a dilemma: inverse-frequency weighting is parameter-free but explodes, whereas more stable variants regain control only by adding sensitive hyperparameters (Lin et al., 2017; Cui et al., 2019). Logarithmic compression offers a middle ground that, to our knowledge, has not been systematically studied as a re-weighting principle across application domains. LWCE realizes it without any additional hyperparameter, PLWCE adds a single exponent that tunes the focusing strength, and ES-LWCE adds an  $\epsilon$  term that keeps the weight bounded and well defined for the rarest classes. Table 2 contrasts these properties with the main cost-sensitive baselines: WCE is parameter-free but its weights are unbounded, Focal and CB Loss regain stability only at the cost of an extra hyperparameter, whereas LWCE is a frequency-based, parameter-free scheme whose weights stay bounded. We evaluate this family against the representative cost-sensitive losses above under a unified protocol spanning image, tabular, and network-intrusion domains—each of which sustains its own active line of imbalance research, from dedicated tabular benchmarks (Liu, Li, Yang, Wei, Kang, Zhu, Hamann, He and Tong, 2025) to intrusion detection under severe class skew (Abdelkhalek and Mashaly, 2023).

**Table 2**

Positioning of LWCE against representative cost-sensitive losses. “Frequency”/“Difficulty” denote the re-weighting signal; “Param.-free” means no loss hyperparameter is required; “Bounded weight” refers to whether the multiplicative weighting factor is prevented from unbounded growth under extreme class imbalance. For Focal Loss this factor is the per-sample difficulty modulation  $(1 - p_i)^\gamma$ , which is inherently confined to  $[0, 1]$  rather than a class-frequency weight; LWCE instead compresses the class-frequency weight itself. Logit Adjustment uses the class prior but applies it as an *additive* shift to the logits rather than a multiplicative weight, so the bounded-weight criterion does not directly apply (n/a).

| Loss               | Frequency | Difficulty | Param.-free | Bounded weight |
|--------------------|-----------|------------|-------------|----------------|
| WCE                | ✓         | –          | ✓           | –              |
| Focal Loss         | –         | ✓          | –           | ✓              |
| CB Loss            | ✓         | –          | –           | ✓              |
| Logit Adjustment   | ✓         | –          | –           | n/a            |
| <b>LWCE (ours)</b> | ✓         | –          | ✓           | ✓              |

## 3. Proposed Method

We study a single family of re-weighting losses built on logarithmic weight compression. All three members replace the inverse-frequency weight of WCE with a logarithmically compressed counterpart and differ only in how that base weight is subsequently shaped. The three formulations and the distinct role of each are summarized in Table 3 and defined precisely below.

### 3.1. The logarithmic re-weighting family

Consider a classification problem with  $C$  classes, where class  $c$  contains  $n_c$  training samples. All three losses re-weight the standard cross-entropy objective by a class-dependent factor  $w_c$ ,

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N w_{y_i} \log p_{i,y_i}, \quad (2)$$

where  $p_{i,y_i}$  is the predicted probability of the true class  $y_i$  of sample  $i$  and  $N$  is the number of samples; the members differ only in the choice of  $w_c$ . Log-Weighted Cross-Entropy (LWCE) uses the parameter-free weight

$$w_c^{\text{LWCE}} = \frac{1}{\log(1 + n_c)}. \quad (3)$$

Power-Log Weighted Cross-Entropy (PLWCE) raises the logarithmic denominator to a single exponent  $\alpha > 0$ ,

$$w_c^{\text{PLWCE}} = \frac{1}{(\log(1 + n_c))^\alpha}, \quad (4)$$



which continuously tunes how strongly rare classes are emphasized and reduces to LWCE at  $\alpha = 1$ .  $\varepsilon$ -Smoothed Log-Weighted Cross-Entropy (ES-LWCE) adds a small constant  $\varepsilon > 0$  to the denominator,

$$w_c^{\text{ES-LWCE}} = \frac{1}{\log(1 + n_c) + \varepsilon}, \quad (5)$$

which caps the weight at  $1/\varepsilon$ .

*Weight normalization.* In all three formulations the raw weights are mean-normalized before use,

$$\tilde{w}_c = \frac{C w_c}{\sum_{j=1}^C w_j}, \quad (6)$$

so that  $\frac{1}{C} \sum_c \tilde{w}_c = 1$  and the overall loss scale is independent of the weighting scheme; only the *relative* emphasis across classes differs. We apply the identical normalization to every re-weighting baseline, including WCE and Class-Balanced loss, so that all methods share one scale convention and differences in performance cannot be attributed to a global change in effective learning rate. All properties below are stated for the raw weights  $w_c$ ; since Eq. (6) is a single positive rescaling common to every class, it leaves the rare-to-frequent *ratios* of Propositions 3 and 4 unchanged.

**Table 3**

The proposed logarithmic re-weighting family.  $n_c$  denotes the number of training samples in class  $c$ ;  $\alpha > 0$  and  $\varepsilon > 0$  are the only free parameters, belonging to PLWCE and ES-LWCE respectively.

| Loss    | Class weight $w_c$                      | Role                                                 |
|---------|-----------------------------------------|------------------------------------------------------|
| LWCE    | $\frac{1}{\log(1 + n_c)}$               | Parameter-free base weight                           |
| PLWCE   | $\frac{1}{(\log(1 + n_c))^\alpha}$      | Exponent $\alpha$ tunes emphasis strength            |
| ES-LWCE | $\frac{1}{\log(1 + n_c) + \varepsilon}$ | Constant $\varepsilon$ bounds and smooths the weight |

### 3.2. Theoretical properties

We now make precise the sense in which logarithmic compression tames the weight explosion of WCE. Throughout, let  $\rho = n_{\max}/n_{\min} \geq 1$  denote the imbalance ratio between the most and least frequent classes, and let the subscripts min and max index the rarest and the most frequent class, respectively. Four properties characterize the proposed family.

*Proposition 1 (Monotonicity).* For LWCE, PLWCE, and ES-LWCE, the weight  $w_c$  is strictly decreasing in the class count  $n_c$ . Consequently a rarer class always receives a larger weight, which preserves the qualitative behaviour of inverse-frequency weighting. *Proof.*  $\log(1 + n_c)$  is strictly increasing in  $n_c$ , so its reciprocal, any positive power of its reciprocal, and the reciprocal of its  $\varepsilon$ -shift are all strictly decreasing in  $n_c$ .  $\square$

*Proposition 2 (Boundedness).* For  $n_c \geq 1$  the individual weights are bounded:  $0 < w_c^{\text{LWCE}} \leq 1/\log 2$  and  $0 < w_c^{\text{PLWCE}} \leq 1/(\log 2)^\alpha$ . The ES-LWCE weight is bounded even without this assumption:  $0 < w_c^{\text{ES-LWCE}} \leq 1/\varepsilon$  for every  $n_c \geq 0$ . *Proof.* Each weight is maximized at its smallest admissible count by Proposition 1. For LWCE and PLWCE this is  $n_c = 1$ , giving  $1/\log 2$  and  $1/(\log 2)^\alpha$ . For ES-LWCE the denominator is at least  $\varepsilon$ , attained as  $n_c \rightarrow 0$ , so the weight never exceeds  $1/\varepsilon$ .  $\square$

This proposition also clarifies the distinct role of ES-LWCE. LWCE is already bounded for every class that appears at least once; what ES-LWCE adds is robustness at  $n_c = 0$ , where the LWCE denominator  $\log(1 + n_c)$  vanishes and the weight diverges. Such vanishing counts arise, for instance, when class frequencies are estimated per mini-batch. ES-LWCE turns the implicit bound of LWCE into an explicit, user-controlled cap of  $1/\varepsilon$ .


**Proposition 3 (Dynamic-range compression).** Fix the minority count  $n_{\min} \geq 1$  and let  $\rho = n_{\max}/n_{\min} \rightarrow \infty$  through growth of  $n_{\max}$ . The rare-to-frequent weight ratio then grows *linearly* in  $\rho$  under WCE but only *logarithmically* under the proposed family. Under WCE,

$$\frac{w_{\min}^{\text{WCE}}}{w_{\max}^{\text{WCE}}} = \frac{n_{\max}}{n_{\min}} = \rho, \quad (7)$$

whereas under LWCE and PLWCE,

$$\frac{w_{\min}^{\text{LWCE}}}{w_{\max}^{\text{LWCE}}} = \frac{\log(1 + n_{\max})}{\log(1 + n_{\min})} = O(\log \rho), \quad \frac{w_{\min}^{\text{PLWCE}}}{w_{\max}^{\text{PLWCE}}} = \left( \frac{\log(1 + n_{\max})}{\log(1 + n_{\min})} \right)^\alpha = O((\log \rho)^\alpha). \quad (8)$$

Comparing Eqs. (7) and (8), the compression reduces the dynamic range of the weights from linear to logarithmic in  $\rho$  while preserving the ordering that favours rare classes. Note that this is a statement about the *ratio*: unlike the individual weights of Proposition 2, the ratio is not constant-bounded but grows without limit as  $\rho \rightarrow \infty$ —only much more slowly than the linear growth of WCE. Figure 1 makes this compression concrete: with the majority count fixed at  $n_{\max} = 10^4$ , the rarest class receives a weight  $\rho = 10^4$  times larger than the majority under WCE, but only about 13 times larger under LWCE, with PLWCE and ES-LWCE interpolating between these regimes.



**Figure 1:** Dynamic-range compression (Proposition 3). The vertical axis is the rare-to-frequent weight ratio  $w_{\min}/w_{\max}$  as a function of the minority count  $n_c$ , with the majority count fixed at  $n_{\max} = 10^4$ ; this ratio equals the imbalance ratio  $\rho$  for WCE and is invariant to the mean-normalization of Eq. (6). Inverse-frequency (WCE) weighting explodes linearly to  $\rho = 10^4$  as  $n_c \rightarrow 1$ , and the inverse-square-root variant to  $\sqrt{\rho}$ , whereas the logarithmic family stays within a small multiple of the majority baseline (LWCE peaks near 13). ES-LWCE lies at or below LWCE, confirming that its  $\varepsilon$  term only softens the weight rather than amplifying it.

**Proposition 4 (Gradient-scale control).** In weighted cross-entropy the per-sample gradient with respect to the logits is the unweighted cross-entropy gradient scaled by the class weight  $w_{y_i}$  (Eq. (2)). The full gradient magnitude also depends on the predicted probabilities and the feature representation, so we do not claim that the minority-to-majority



gradient ratio equals the weight ratio exactly. Rather, the *class-dependent multiplicative component* of the gradient scale follows the same asymptotic order as the corresponding weight ratio of Proposition 3: it is  $O(\rho)$  under WCE but  $O(\log \rho)$  under LWCE. Logarithmic compression therefore limits the extent to which minority gradients can overwhelm the majority under extreme imbalance, which is the mechanism by which LWCE mitigates the optimization instability of WCE; ES-LWCE additionally caps the worst-case class-dependent factor at  $1/\varepsilon$ . Section 4 tests this mechanism directly by tracking the per-epoch gradient norm and the empirical minority-to-majority gradient ratio.

### 3.3. Parameter selection for PLWCE

LWCE and ES-LWCE require essentially no tuning: LWCE is parameter-free, and the  $\varepsilon$  of ES-LWCE is a user-set cap rather than an accuracy knob (Section 3.4). PLWCE, however, exposes a genuine hyperparameter  $\alpha$  that controls the focusing strength, with three qualitatively distinct regimes:

- $0 < \alpha < 1$ : *mild* re-weighting, weaker than LWCE;
- $\alpha = 1$ : exactly LWCE;
- $\alpha > 1$ : *strong* tail emphasis.

We select  $\alpha$  with a reproducible, data-driven procedure that never touches the test set. For each dataset (and each imbalance level, where applicable) we draw a stratified *proxy* subset from the *training* split—a fixed fraction of the training data, with a minimum number of samples per class so that rare classes are never eliminated from the proxy—and hold out a validation portion of that proxy. A model is trained on the proxy for a reduced number of epochs at each candidate  $\alpha$  on a fixed grid, and we retain the value that maximizes the macro-averaged  $F_1$  on the proxy validation set. Only this single scalar is then used to train the final model under the full protocol. The search is therefore confined to the training split, is run once per dataset, and adds a bounded, one-time cost (the number of grid points times the reduced-epoch proxy budget) on top of a single full run; it introduces no per-iteration overhead and no test-set leakage.

### 3.4. ES-LWCE and practical use cases

ES-LWCE is not intended to outperform LWCE on datasets whose class counts are fixed and strictly positive; there, LWCE is already bounded (Proposition 2) and the two coincide as  $\varepsilon \rightarrow 0$ . Its value appears when the class counts used to form the weights can *vanish or drift*. When frequencies are estimated per mini-batch, or maintained online for streaming or continually growing label sets, a class may be momentarily absent ( $n_c = 0$ ); the LWCE weight then diverges, whereas ES-LWCE remains well defined with weight at most  $1/\varepsilon$ . The constant  $\varepsilon$  trades off two effects:

- small  $\varepsilon$ : stronger minority emphasis, closer to LWCE;
- large  $\varepsilon$ : more conservative gradient scaling, closer to uniform.

Because  $\varepsilon$  shifts the weights monotonically from the LWCE limit toward the uniform (Cross-Entropy) limit and never past LWCE, it acts purely as a smoothing/robustness control. We therefore position ES-LWCE as an extension for vanishing- or batch-adaptive-count settings rather than as a method that is required at the dataset level.

### 3.5. Computational complexity

All three losses add negligible cost over plain Cross-Entropy. The class weights depend only on the class counts and are precomputed once from the training split; there are no additional model parameters and no extra backward-pass terms. Once the weights are precomputed, LWCE and ES-LWCE incur the same asymptotic training complexity as WCE, requiring only a scalar multiplication per sample. The only added cost specific to PLWCE is the one-time grid search of Section 3.3, which is bounded and performed on a reduced proxy rather than at full training scale.

## 4. Experiments

Our experiments are organized around four questions:

1. Does the advantage of the logarithmic family grow as the imbalance becomes more severe?
2. Does LWCE empirically reduce the optimization instability of WCE?
3. Does the PLWCE exponent  $\alpha$  control tail emphasis as intended?

#### 4. Is ES-LWCE stable under vanishing or extremely rare class counts?

To probe architecture- and modality-independence, we evaluate the same loss family across image classification (CIFAR-10/100-LT with a ResNet-32 backbone), tabular classification and network-intrusion detection (multilayer perceptrons), and text classification (a compact Transformer encoder on a long-tailed emotion benchmark).

### 4.1. Baselines and evaluation

We compare the three proposed losses (LWCE, PLWCE, ES-LWCE) against Cross-Entropy, Weighted Cross-Entropy, an inverse-square-root variant, Class-Balanced loss (Cui et al., 2019), Focal loss (Lin et al., 2017), and Logit Adjustment (Menon et al., 2021) as a margin/posterior-shift baseline. Because overall accuracy is uninformative under strong imbalance, we report macro-averaged  $F_1$  and worst-class accuracy as the primary metrics, and additionally balanced accuracy, the geometric mean of per-class recall, minority-class recall, and the full per-class accuracy vector. All results are averaged over five random seeds and reported as mean  $\pm$  standard deviation.

### 4.2. Hyperparameter selection

A recurring concern with re-weighting losses is that any apparent advantage may stem from tuning rather than from the weighting scheme itself. To rule this out, every baseline that exposes a free parameter is tuned under an identical, automated budget, while LWCE is used exactly as defined in Eq. (3) with *no* search at all. Concretely, the PLWCE and inverse-power exponents, the Focal focusing parameter, the ES-LWCE smoothing constant, and the Logit-Adjustment temperature are each selected on a validation split by the procedure of Section 3.3; Cross-Entropy, WCE, the inverse-square-root variant, and Class-Balanced loss are parameter-free or use their authors' recommended value. Search uses a deterministic grid over each parameter (with the Focal grid lower-bounded at  $\gamma = 1$ , below which the modulating factor  $(1 - p_t)^\gamma$  yields an unstable gradient as  $p_t \rightarrow 1$ ), so that the entire admissible range is covered reproducibly, and the smoothing constant  $\epsilon$  is searched on a logarithmic grid because its effect on the weights is multiplicative. LWCE therefore competes with no tuning against baselines that each receive a full search.

### 4.3. Main results

### 4.4. Optimization stability and mechanism

### 4.5. Ablation and sensitivity

We isolate the contribution of each ingredient of the logarithmic family—the logarithmic compression itself, the power exponent  $\alpha$ , and the smoothing constant  $\epsilon$ —through the ablation of Table 4. Reading from WCE (no compression) to LWCE isolates the effect of logarithmic compression; adding  $\alpha$  gives PLWCE and adding  $\epsilon$  gives ES-LWCE; the *Combined* variant activates both extensions at once,  $w_c \propto (\log(1 + n_c) + \epsilon)^{-\alpha}$ , and is included—following the same design as the other rows—to probe whether the two extensions are worth merging into a single loss.

We do not propose the Combined variant as a method. Across all domains it yields no improvement over PLWCE alone: the selection procedure of Section 3.3 consistently drives  $\epsilon$  to the lower end of its grid, so the Combined weight collapses back onto PLWCE. This is a useful negative result rather than a failure—it shows that, on datasets with fixed and strictly positive class counts, the smoothing dimension is inactive for accuracy, exactly as anticipated in Section 3.4:  $\epsilon$  is a robustness control for vanishing or drifting counts, not an accuracy lever. It therefore supports treating PLWCE and ES-LWCE as separate tools rather than a single combined loss.

**Table 4**

Ablation of the logarithmic family by ingredient: logarithmic compression, the power exponent  $\alpha$ , and the smoothing constant  $\epsilon$ . The Combined variant activates both extensions and is reported only to test whether merging them helps; it is not proposed as a method. (Metrics averaged over five seeds; full per-domain numbers in the supplementary material.)

| Method   | Log compression | Power $\alpha$ | Smoothing $\epsilon$ | Macro- $F_1$ |
|----------|-----------------|----------------|----------------------|--------------|
| WCE      | —               | —              | —                    |              |
| LWCE     | ✓               | —              | —                    |              |
| PLWCE    | ✓               | ✓              | —                    |              |
| ES-LWCE  | ✓               | —              | ✓                    |              |
| Combined | ✓               | ✓              | ✓                    |              |

## 5. Conclusion

### CRedit authorship contribution statement

**Seung Ho Go:** Conceptualization, Methodology, Software, Writing – original draft. **Hyunseok Lee:** Methodology, Validation, Writing – review & editing. **Jin Hee Yoon:** Supervision, Conceptualization, Writing – review & editing.

## References

- Abdelkhalik, A., Mashaly, M., 2023. Addressing the class imbalance problem in network intrusion detection systems using data resampling and deep learning. *The Journal of Supercomputing* 79, 10611–10644.
- Branco, P., Torgo, L., Ribeiro, R.P., 2016. A survey of predictive modeling on imbalanced domains. *ACM Computing Surveys* 49, 1–50.
- Buda, M., Maki, A., Mazurowski, M.A., 2018. A systematic study of the class imbalance problem in convolutional neural networks. *Neural Networks* 106, 249–259.
- Byrd, J., Lipton, Z.C., 2019. What is the effect of importance weighting in deep learning?, in: *ICML*.
- Cai, J., Wang, Y., Hwang, J.N., 2021. ACE: Ally complementary experts for solving long-tailed recognition in one-shot, in: *ICCV*.
- Cao, K., Wei, C., Gaidon, A., Aréchiga, N., Ma, T., 2019. Learning imbalanced datasets with label-distribution-aware margin loss, in: *NeurIPS*.
- Chawla, N.V., Bowyer, K.W., Hall, L.O., Kegelmeyer, W.P., 2002. SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research* 16, 321–357.
- Chen, W., Yang, K., Yu, Z., Shi, Y., Chen, C.L.P., 2024. A survey on imbalanced learning: Latest research, applications and future directions. *Artificial Intelligence Review* 57, 137.
- Cui, Y., Jia, M., Lin, T.Y., Song, Y., Belongie, S., 2019. Class-balanced loss based on effective number of samples, in: *CVPR*.
- Drummond, C., Holte, R.C., 2003. C4.5, class imbalance, and cost sensitivity: Why under-sampling beats over-sampling, in: *Workshop on Learning from Imbalanced Datasets, ICML*.
- Guo, D., Li, Z., Zheng, M., Zhao, H., Zhou, M., Zha, H., 2022. Learning to re-weight examples with optimal transport for imbalanced classification, in: *NeurIPS*.
- He, H., Bai, Y., Garcia, E.A., Li, S., 2008. ADASYN: Adaptive synthetic sampling approach for imbalanced learning, in: *International Joint Conference on Neural Networks (IJCNN)*, pp. 1322–1328.
- He, H., Garcia, E.A., 2009. Learning from imbalanced data. *IEEE Transactions on Knowledge and Data Engineering* 21, 1263–1284.
- Johnson, J.M., Khoshgoftaar, T.M., 2019. Survey on deep learning with class imbalance. *Journal of Big Data* 6, 27.
- Kang, B., Xie, S., Rohrbach, M., Yan, Z., Gordo, A., Feng, J., Kalantidis, Y., 2020. Decoupling representation and classifier for long-tailed recognition, in: *ICLR*.
- Khan, S.H., Hayat, M., Bennamoun, M., Soheli, F.A., Togneri, R., 2018. Cost-sensitive learning of deep feature representations from imbalanced data. *IEEE Transactions on Neural Networks and Learning Systems* 29, 3573–3587.
- Krawczyk, B., 2016. Learning from imbalanced data: Open challenges and future directions. *Progress in Artificial Intelligence* 5, 221–232.
- Li, B., Liu, Y., Wang, X., 2019. Gradient harmonized single-stage detector, in: *AAAI*.
- Lin, T.Y., Goyal, P., Girshick, R., He, K., Dollár, P., 2017. Focal loss for dense object detection, in: *ICCV*.
- Liu, Z., Li, Z., Yang, Z., Wei, T., Kang, J., Zhu, Y., Hamann, H., He, J., Tong, H., 2025. CLIMB: Class-imbalanced learning benchmark on tabular data, in: *NeurIPS Datasets and Benchmarks Track*.
- Liu, Z., Miao, Z., Zhan, X., Wang, J., Gong, B., Yu, S.X., 2019. Large-scale long-tailed recognition in an open world, in: *CVPR*.
- Menon, A.K., Jayasumana, S., Rawat, A.S., Jain, H., Veit, A., Kumar, S., 2021. Long-tail learning via logit adjustment, in: *ICLR*.
- Ren, J., Yu, C., Sheng, S., Ma, X., Zhao, H., Yi, S., Li, H., 2020. Balanced meta-softmax for long-tailed visual recognition, in: *NeurIPS*.
- Ren, M., Zeng, W., Yang, B., Urtasun, R., 2018. Learning to reweight examples for robust deep learning, in: *ICML*.
- Shi, J.X., Wei, T., Zhou, Z., Shao, J.J., Han, X.Y., Li, Y.F., 2024. Long-tail learning with foundation model: Heavy fine-tuning hurts, in: *ICML*.
- Tan, J., Wang, C., Li, B., Li, Q., Ouyang, W., Yan, C., Yan, J., 2020. Equalization loss for long-tailed object recognition, in: *CVPR*.
- Wang, S., Liu, W., Wu, L., Cao, L., Meng, Q., Chen, W., 2019. Learning from imbalanced data sets with weighted cross-entropy function. *Pattern Recognition Letters* .
- Wang, X., Lian, L., Miao, Z., Liu, Z., Yu, S.X., 2021. Long-tailed recognition by routing diverse distribution-aware experts, in: *ICLR*.
- Xiang, L., Ding, G., Han, J., 2020. Learning from multiple experts: Self-paced knowledge distillation for long-tailed classification, in: *ECCV*.
- Zhang, S., Li, Z., Yan, S., He, X., Sun, J., 2021. Distribution alignment: A unified framework for long-tail visual recognition, in: *CVPR*.
- Zhang, Y., Hooi, B., Hong, L., Feng, J., 2022. Self-supervised aggregation of diverse experts for test-agnostic long-tailed recognition, in: *NeurIPS*.
- Zhang, Y., Kang, B., Hooi, B., Yan, S., Feng, J., 2023. Deep long-tailed learning: A survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 45, 10795–10816.
- Zhou, B., Cui, Q., Wei, X.S., Chen, Z.M., 2020. BBN: Bilateral-branch network with cumulative learning for long-tailed visual recognition, in: *CVPR*.